

Regression Discontinuity

Regression Discontinuity

```
# Install packages
if (!require("pacman")) install.packages("pacman")

## Loading required package: pacman

pacman::p_load(# Tidyverse packages including dplyr and ggplot2
              tidyverse,
              rdd,
              ggthemes,
              tidymodels)

set.seed(1)
```

Definition

In social sciences, a regression discontinuity design is a quasi-experimental pretest-posttest design that elicits the causal effects of interventions by assigning a cutoff or threshold above or below which an intervention is assigned. By comparing observations lying closely on either side of the threshold, it is possible to estimate the average treatment effect in environments in which randomization is unfeasible.

Treatment Using a Running Variable

In an ideal experiment, we would be able to randomly assign our units to treatment and control. However, as we've seen, this is not always possible in social science contexts. Let's consider a classic question from political science: do incumbent politicians enjoy an incumbency advantage? In other words, do incumbents garner a higher vote share than they would if they were running for the first time?

To explore this question, we are going to use data taken from Lee (2008). We have a few variables to define:

- **difshare**: Normalized proportion of vote share the party received in the last election
- **yeare1**: Year of current election
- **myoutcomenext**: 0/1 binary for whether the candidate won re-election
- **win_relection**: "win"/"lose" whether the candidate won re-election
- **incumbent**: 0/1 for whether the candidate is an incumbent

```
elections <- read_csv('../data/indiv_final.csv')
```


```
## Parsed with column specification:
## cols(
##   yearel = col_double(),
##   myoutcomenext = col_double(),
##   difshare = col_double()
## )
```

```
elections <- elections %>%
  mutate(win_reelection = ifelse(myoutcomenext == 1,
                                'win',
                                'lose')) %>%
  mutate(incumbent = ifelse(difshare > 0,
                            1,
                            0))
```

Suppose we are interested in the effect of incumbency on the probability of winning re-election. We might look at the distribution of vote shares by incumbency. Let's look at this boxplot:

```
elections %>%
  mutate(incumbent = as.factor(incumbent)) %>%
  ggplot() +
  geom_boxplot(aes(x = incumbent,
                  y = difshare,
                  group = incumbent,
                  fill = incumbent)) +
  ggtitle('Boxplot of Vote Share Among \n Incumbents') +
  theme_fivethirtyeight()
```

Boxplot of Vote Share Among Incumbents



Looks like incumbents enjoy a pretty significant advantage! Let's investigate this further by this using a linear probability model:


```
election_lm <- lm(myoutcomenext ~ incumbent,
                  data = elections)
summary(election_lm)
```

```
##
## Call:
## lm(formula = myoutcomenext ~ incumbent, data = elections)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.78141 -0.01401 -0.01401  0.21859  0.98599
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.014012   0.002259   6.203 5.61e-10 ***
## incumbent    0.767395   0.003576 214.604 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2765 on 24935 degrees of freedom
## Multiple R-squared:  0.6488, Adjusted R-squared:  0.6487
## F-statistic: 4.606e+04 on 1 and 24935 DF, p-value: < 2.2e-16
```

Incumbency is a statistically significant covariate!

What might be the problem with using .77 as the coefficient in this case? Let's take at the distribution of previous elections voteshares between winners and losers.

```
elections %>%
  mutate(incumbent = as.factor(incumbent)) %>%
  mutate(decile = as.factor(ntile(difshare, 10))) %>%
  ggplot() +
  geom_bar(aes(y = decile,
               fill = win_reelection),
           position = 'dodge') +
  ggtitle('Barplot of Re-Election Winners by \n Previous Election Vote Share Decile') +
  theme_fivethirtyeight() +
  coord_flip()
```



Question: Groups 1 - 4 look ok - by definition someone cannot win reelection if they lost the last election. But what about the difference in distributions between deciles like 5 and 6, versus the distributions in deciles 7-10? Why would these different distributions pose a problem for trusting our previous point estimate?


Solution: Our main worry is that people who won huge proportions of the vote in the last election are systematically different from those who barely won in the last election. For instance, it should not be surprising when a Democratic incumbent carries CA-13 (Berkeley's congressional district) not because they enjoy an incumbency advantage but because the district is heavily Democratic for other reasons.

Running Variable

We might assume that the selection into treatment/control conditions is driven entirely by observable characteristics (selection on the observables). If this was the case, then we could add these characteristics as controls to our regression, and we would be ok. In practice, this is rarely a realistic assumption though, and we most likely will worry about selection on unobservable characteristics - essentially confounders that we do not see but affect both treatment and the outcome.

The basic intuition behind regression discontinuity designs is that we use a **running variable** that determines whether a unit was assigned to either treatment or control. Let's take a look:

```
elections %>%
  ggplot() +
  geom_density(aes(x = difshare)) +
  ggtitle('Density of Normalized Voteshare') +
  geom_vline(xintercept = 0) +
  theme_fivethirtyeight()
```



“0” here is the cutpoint we use to assign someone to either incumbent or non-incumbent treatment conditions. The basic logic behind the RD is that those individuals on either side of the cutpoints will be very similar to each other in terms of baseline covariates (on both observed and unobserved characteristics).

McCrary Density Test


Before we estimate model for individuals on either side of the cutpoint though, we might be concerned about their manipulation into treatment and control. For example, if the running variable was a passing test score

to move onto to the next grade, you might imagine that a teacher bumps up a student from a 59 to a 60. Similarly, if the cutoff to be recruited into an honors program is a 90, you might worry that a student with an 89 who knows that they could appeal to be admitted anyway differs from the student who does not think such a thing is negotiable.

McCrary proposes a test for this kind of problem. Specifically, he motivates the test by giving an example of patients who are assigned to either waiting room A or B, but only waiting room A receives the experimental treatment. Patients learn about this fact, so as those who are assigned to waiting room B are walking there, they instead decide to go to waiting room A. Given enough patients doing this, we should expect waiting room A to become crowded and waiting room B to be relatively empty. Let's see what that looks like graphically in our dataset:

```
elections %>%
  filter(difshare > -.05 & difshare < .05) %>%
  ggplot() +
  geom_histogram(aes(x = difshare,
                     group = as.factor(incumbent),
                     fill = as.factor(incumbent)))
```


```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



It looks like there isn't evidence of sorting at the cutpoint - in fact the distributions look identical. McCrary points out that these types of histograms can be biased at the cutpoint. He instead advocates for using local linear regressions to smooth the histograms at the cutpoint. Luckily, this procedure is implemented for us in the rdd package:

```
DCdensity(elections$difshare,
          cutpoint = 0,
          verbose = TRUE,
          ext.out = TRUE,
          htest = TRUE)
```

```
## Using calculated bin size: 0.005
## Using calculated bandwidth: 0.183
```



```
## Log difference in heights is -0.002 with SE 0.052
## this gives a z-stat of -0.048
## and a p value of 0.962
```

```
## $theta
## [1] -0.002470001
##
## $se
## [1] 0.05193608
##
## $z
## [1] -0.04755847
##
## $p
## [1] 0.9620681
```

```

##
## $binsize
## [1] 0.005290885
##
## $bw
## [1] 0.1833466
##
## $cutpoint
## [1] 0
##
## $data
##           cellmp      cellval
## 1  -1.002622783 0.01515854
## 2  -0.997331898 1.60680547
## 3  -0.992041013 0.21979886
## 4  -0.986750127 0.07579271
## 5  -0.981459242 0.03789636
## 6  -0.976168356 0.12884761
## 7  -0.970877471 0.12126834
## 8  -0.965586586 0.07579271
## 9  -0.960295700 0.11368907
## 10 -0.955004815 0.08337198
## 11 -0.949713929 0.06821344
## 12 -0.944423044 0.09853052
## 13 -0.939132159 0.20464032
## 14 -0.933841273 0.08337198
## 15 -0.928550388 0.16674396
## 16 -0.923259502 0.15916469
## 17 -0.917968617 0.07579271
## 18 -0.912677732 0.06063417
## 19 -0.907386846 0.11368907
## 20 -0.902095961 0.12884761
## 21 -0.896805075 0.12126834
## 22 -0.891514190 0.06063417
## 23 -0.886223305 0.18948178
## 24 -0.880932419 0.21221959
## 25 -0.875641534 0.16674396
## 26 -0.870350648 0.09853052
## 27 -0.865059763 0.11368907
## 28 -0.859768878 0.17432323
## 29 -0.854477992 0.21221959
## 30 -0.849187107 0.15158542
## 31 -0.843896221 0.09853052
## 32 -0.838605336 0.18948178
## 33 -0.833314451 0.18948178
## 34 -0.828023565 0.11368907
## 35 -0.822732680 0.12126834
## 36 -0.817441794 0.26527449
## 37 -0.812150909 0.12884761
## 38 -0.806860024 0.18948178
## 39 -0.801569138 0.15158542
## 40 -0.796278253 0.18190251
## 41 -0.790987367 0.19706105
## 42 -0.785696482 0.21221959

```



```

## 43 -0.780405597 0.20464032
## 44 -0.775114711 0.24253667
## 45 -0.769823826 0.18948178
## 46 -0.764532940 0.19706105
## 47 -0.759242055 0.26527449
## 48 -0.753951170 0.32590866
## 49 -0.748660284 0.29559157
## 50 -0.743369399 0.38654282
## 51 -0.738078513 0.30317084
## 52 -0.732787628 0.33348793
## 53 -0.727496743 0.33348793
## 54 -0.722205857 0.47749408
## 55 -0.716914972 0.39412210
## 56 -0.711624086 0.45475626
## 57 -0.706333201 0.46991481
## 58 -0.701042316 0.53812825
## 59 -0.695751430 0.68971367
## 60 -0.690460545 0.55328679
## 61 -0.685169659 0.52296970
## 62 -0.679878774 0.60634169
## 63 -0.674587889 0.59118314
## 64 -0.669297003 0.56086606
## 65 -0.664006118 0.78066492
## 66 -0.658715232 0.81098200
## 67 -0.653424347 0.81856128
## 68 -0.648133462 0.79582346
## 69 -0.642842576 0.83371982
## 70 -0.637551691 0.88677472
## 71 -0.632260805 0.97014670
## 72 -0.626969920 0.74276857
## 73 -0.621679035 0.81856128
## 74 -0.616388149 0.73518929
## 75 -0.611097264 0.92467107
## 76 -0.605806378 0.94740888
## 77 -0.600515493 0.91709180
## 78 -0.595224608 0.77308565
## 79 -0.589933722 0.86403690
## 80 -0.584642837 0.81098200
## 81 -0.579351951 0.81856128
## 82 -0.574061066 0.99288451
## 83 -0.568770181 0.87161617
## 84 -0.563479295 0.93225034
## 85 -0.558188410 0.99288451
## 86 -0.552897524 0.85645763
## 87 -0.547606639 0.88677472
## 88 -0.542315754 0.76550638
## 89 -0.537024868 0.84129909
## 90 -0.531733983 0.96256743
## 91 -0.526443097 0.88677472
## 92 -0.521152212 1.34153098
## 93 -0.515861327 0.94740888
## 94 -0.510570441 0.95498816
## 95 -0.505279556 1.04593941
## 96 -0.499988670 1.12173212

```

97 -0.494697785 1.13689066
98 -0.489406900 0.94740888
99 -0.484116014 1.05351868
100 -0.478825129 0.73518929
101 -0.473534243 0.78824419
102 -0.468243358 0.78824419
103 -0.462952473 0.65939658
104 -0.457661587 0.65939658
105 -0.452370702 0.68213440
106 -0.447079816 0.62907950
107 -0.441788931 0.59118314
108 -0.436498046 0.62150023
109 -0.431207160 0.75034784
110 -0.425916275 0.57602460
111 -0.420625389 0.57602460
112 -0.415334504 0.64423804
113 -0.410043619 0.62907950
114 -0.404752733 0.59876241
115 -0.399461848 0.59876241
116 -0.394170962 0.62150023
117 -0.388880077 0.53812825
118 -0.383589192 0.48507335
119 -0.378298306 0.54570752
120 -0.373007421 0.56844533
121 -0.367716535 0.64423804
122 -0.362425650 0.63665877
123 -0.357134765 0.68213440
124 -0.351843879 0.57602460
125 -0.346552994 0.77308565
126 -0.341262108 0.66697585
127 -0.335971223 0.63665877
128 -0.330680338 0.74276857
129 -0.325389452 0.79582346
130 -0.320098567 0.71245148
131 -0.314807681 0.72761002
132 -0.309516796 0.67455513
133 -0.304225911 0.65181731
134 -0.298935025 0.86403690
135 -0.293644140 0.72761002
136 -0.288353254 0.72003075
137 -0.283062369 0.71245148
138 -0.277771484 0.78824419
139 -0.272480598 0.74276857
140 -0.267189713 0.70487221
141 -0.261898827 0.61392096
142 -0.256607942 0.74276857
143 -0.251317057 0.83371982
144 -0.246026171 0.59118314
145 -0.240735286 0.75792711
146 -0.235444400 0.84887836
147 -0.230153515 0.76550638
148 -0.224862630 0.78066492
149 -0.219571744 0.78824419
150 -0.214280859 0.74276857

151 -0.208989973 0.79582346
152 -0.203699088 0.64423804
153 -0.198408203 0.57602460
154 -0.193117317 0.74276857
155 -0.187826432 0.75034784
156 -0.182535546 0.93225034
157 -0.177244661 0.77308565
158 -0.171953776 0.65939658
159 -0.166662890 0.68213440
160 -0.161372005 0.75792711
161 -0.156081119 0.70487221
162 -0.150790234 0.76550638
163 -0.145499349 0.72003075
164 -0.140208463 0.82614055
165 -0.134917578 0.89435399
166 -0.129626692 0.77308565
167 -0.124335807 0.72003075
168 -0.119044922 0.75034784
169 -0.113754036 0.84129909
170 -0.108463151 0.75034784
171 -0.103172265 0.72003075
172 -0.097881380 0.73518929
173 -0.092590495 0.90193326
174 -0.087299609 0.72761002
175 -0.082008724 0.68213440
176 -0.076717838 0.78824419
177 -0.071426953 0.74276857
178 -0.066136068 0.80340273
179 -0.060845182 0.87919544
180 -0.055554297 0.80340273
181 -0.050263411 0.81098200
182 -0.044972526 0.82614055
183 -0.039681641 1.04593941
184 -0.034390755 0.76550638
185 -0.029099870 0.78824419
186 -0.023808984 0.76550638
187 -0.018518099 0.74276857
188 -0.013227214 0.77308565
189 -0.007936328 0.68971367
190 -0.002645443 0.70487221
191 0.002645443 0.70487221
192 0.007936328 0.68971367
193 0.013227214 0.77308565
194 0.018518099 0.74276857
195 0.023808984 0.76550638
196 0.029099870 0.78824419
197 0.034390755 0.75792711
198 0.039681641 1.04593941
199 0.044972526 0.82614055
200 0.050263411 0.81098200
201 0.055554297 0.77308565
202 0.060845182 0.87919544
203 0.066136068 0.79582346
204 0.071426953 0.74276857

205 0.076717838 0.78066492
206 0.082008724 0.68213440
207 0.087299609 0.72761002
208 0.092590495 0.90193326
209 0.097881380 0.72761002
210 0.103172265 0.72003075
211 0.108463151 0.74276857
212 0.113754036 0.84129909
213 0.119044922 0.75034784
214 0.124335807 0.72003075
215 0.129626692 0.77308565
216 0.134917578 0.89435399
217 0.140208463 0.82614055
218 0.145499349 0.72003075
219 0.150790234 0.75034784
220 0.156081119 0.70487221
221 0.161372005 0.75792711
222 0.166662890 0.67455513
223 0.171953776 0.65939658
224 0.177244661 0.77308565
225 0.182535546 0.92467107
226 0.187826432 0.75034784
227 0.193117317 0.73518929
228 0.198408203 0.57602460
229 0.203699088 0.64423804
230 0.208989973 0.78066492
231 0.214280859 0.72761002
232 0.219571744 0.78824419
233 0.224862630 0.76550638
234 0.230153515 0.75792711
235 0.235444400 0.84129909
236 0.240735286 0.75792711
237 0.246026171 0.59118314
238 0.251317057 0.81098200
239 0.256607942 0.72761002
240 0.261898827 0.59118314
241 0.267189713 0.70487221
242 0.272480598 0.73518929
243 0.277771484 0.76550638
244 0.283062369 0.69729294
245 0.288353254 0.71245148
246 0.293644140 0.72003075
247 0.298935025 0.85645763
248 0.304225911 0.63665877
249 0.309516796 0.65939658
250 0.314807681 0.72003075
251 0.320098567 0.70487221
252 0.325389452 0.77308565
253 0.330680338 0.71245148
254 0.335971223 0.62150023
255 0.341262108 0.65181731
256 0.346552994 0.75034784
257 0.351843879 0.57602460
258 0.357134765 0.65181731

259 0.362425650 0.60634169
260 0.367716535 0.62150023
261 0.373007421 0.53812825
262 0.378298306 0.52296970
263 0.383589192 0.46233554
264 0.388880077 0.48507335
265 0.394170962 0.56844533
266 0.399461848 0.53812825
267 0.404752733 0.53054898
268 0.410043619 0.54570752
269 0.415334504 0.55328679
270 0.420625389 0.50781116
271 0.425916275 0.50023189
272 0.431207160 0.62907950
273 0.436498046 0.53054898
274 0.441788931 0.46991481
275 0.447079816 0.49265262
276 0.452370702 0.50023189
277 0.457661587 0.44717699
278 0.462952473 0.45475626
279 0.468243358 0.48507335
280 0.473534243 0.45475626
281 0.478825129 0.34864647
282 0.484116014 0.53054898
283 0.489406900 0.44717699
284 0.494697785 0.33348793
285 0.499988670 0.42443918
286 0.505279556 0.36380501
287 0.510570441 0.20464032
288 0.515861327 0.21221959
289 0.521152212 0.42443918
290 0.526443097 0.15158542
291 0.531733983 0.24253667
292 0.537024868 0.22737813
293 0.542315754 0.24253667
294 0.547606639 0.27285376
295 0.552897524 0.21979886
296 0.558188410 0.18948178
297 0.563479295 0.23495740
298 0.568770181 0.18190251
299 0.574061066 0.21221959
300 0.579351951 0.21221959
301 0.584642837 0.19706105
302 0.589933722 0.14400615
303 0.595224608 0.14400615
304 0.600515493 0.21221959
305 0.605806378 0.16674396
306 0.611097264 0.14400615
307 0.616388149 0.10610980
308 0.621679035 0.12126834
309 0.626969920 0.12126834
310 0.632260805 0.14400615
311 0.637551691 0.14400615
312 0.642842576 0.15916469

```

## 313 0.648133462 0.12884761
## 314 0.653424347 0.11368907
## 315 0.658715232 0.14400615
## 316 0.664006118 0.10610980
## 317 0.669297003 0.11368907
## 318 0.674587889 0.15916469
## 319 0.679878774 0.13642688
## 320 0.685169659 0.06821344
## 321 0.690460545 0.06063417
## 322 0.695751430 0.15158542
## 323 0.701042316 0.14400615
## 324 0.706333201 0.07579271
## 325 0.711624086 0.10610980
## 326 0.716914972 0.09095125
## 327 0.722205857 0.11368907
## 328 0.727496743 0.06821344
## 329 0.732787628 0.07579271
## 330 0.738078513 0.04547563
## 331 0.743369399 0.16674396
## 332 0.748660284 0.06821344
## 333 0.753951170 0.10610980
## 334 0.759242055 0.09095125
## 335 0.764532940 0.08337198
## 336 0.769823826 0.06821344
## 337 0.775114711 0.07579271
## 338 0.780405597 0.07579271
## 339 0.785696482 0.06821344
## 340 0.790987367 0.06063417
## 341 0.796278253 0.07579271
## 342 0.801569138 0.05305490
## 343 0.806860024 0.06821344
## 344 0.812150909 0.05305490
## 345 0.817441794 0.09853052
## 346 0.822732680 0.06821344
## 347 0.828023565 0.03789636
## 348 0.833314451 0.07579271
## 349 0.838605336 0.07579271
## 350 0.843896221 0.03789636
## 351 0.849187107 0.06063417
## 352 0.854477992 0.08337198
## 353 0.859768878 0.09095125
## 354 0.865059763 0.06063417
## 355 0.870350648 0.04547563
## 356 0.875641534 0.10610980
## 357 0.880932419 0.09095125
## 358 0.886223305 0.06821344
## 359 0.891514190 0.02273781
## 360 0.896805075 0.08337198
## 361 0.902095961 0.06063417
## 362 0.907386846 0.04547563
## 363 0.912677732 0.03789636
## 364 0.917968617 0.04547563
## 365 0.923259502 0.08337198
## 366 0.928550388 0.07579271

```

```
## 367 0.933841273 0.07579271
## 368 0.939132159 0.04547563
## 369 0.944423044 0.09095125
## 370 0.949713929 0.04547563
## 371 0.955004815 0.04547563
## 372 0.960295700 0.05305490
## 373 0.965586586 0.03789636
## 374 0.970877471 0.09095125
## 375 0.976168356 0.03031708
## 376 0.981459242 0.02273781
## 377 0.986750127 0.04547563
## 378 0.992041013 0.11368907
## 379 0.997331898 0.65939658
## 380 1.002622783 0.01515854
```

The hypothesis test is looking to see whether the density of the points is statistically different at the cutpoint. Here we do not see evidence of this (it would be hard to manipulate winning an election!).

Sharp Discontinuity

We can go ahead and estimate our model now! Once again the `rdd` package provides a nice function to let us do this:

```
sharp_rdd_model <- RDestimate(# formula. outcome regressed on running variable x
  # add a | c1 + c2... to add additional covariates
  myoutcomenext ~ difshare,
  # data
  data = elections,
  # specify cutpoint
  cutpoint = 0,
  # specify bandwidth from Lee paper
  bw = .25,
  # return a model object
  model = TRUE)

sharp_rdd_model
```

```
##
## Call:
## RDestimate(formula = myoutcomenext ~ difshare, data = elections,
##   cutpoint = 0, bw = 0.25, model = TRUE)
##
## Coefficients:
##      LATE      Half-BW  Double-BW
##    0.5158      0.4725      0.5855
```

Question: How does our LATE compare to the .77 estimate we saw before?

Solution: Much lower, a difference of about 20% in predicted probability of being re-elected.

Fuzzy Discontinuity

We can also easily implement a fuzzy RD design. As we discussed in lecture, a fuzzy RD does not use one cutoff to assign to treatment and control, but rather uses the running variable as an instrument. To do a fuzzy RD, you simply need to add a `z` to your formula to indicate the treatment variable.

```
fuzzy_rdd_model <- RDestimate(# formula. outcome regressed on running variable x  
  # To change this to a fuzzy rdd, add a +z to the formula to indicate the treatment variable  
  # add a | c1 + c2... to add additional covariates  
  myoutcomenext ~ difshare + incumbent,  
  # data  
  data = elections,  
  # specify cutpoint  
  cutpoint = 0,  
  # specify bandwidth from Lee paper  
  bw = .25,  
  # return a model object  
  model = TRUE)
```

```
fuzzy_rdd_model
```

```
##  
## Call:  
## RDestimate(formula = myoutcomenext ~ difshare + incumbent, data = elections,  
##   cutpoint = 0, bw = 0.25, model = TRUE)  
##  
## Coefficients:  
##      LATE      Half-BW    Double-BW  
##    0.5158      0.4725      0.5855
```

Question: Our LATE for both the Sharp and Fuzzy models was the same here. Does that make sense?

Solution: Yes! Because we defined our treatment indicator in sharp terms relative to our running variable, the sharp and fuzzy designs should be identical in this case.

Bandwidth Selection

One of the main drawbacks of the regression discontinuity design is determining the optimal choice of bandwidth around the cutpoint. The intuition is that we want to pick a bandwidth such that the units on either side are very similar on both observed and unobserved characteristics - but if we knew how to do that then we could just use all of the data and matching! One way to select the bandwidth might be theory-driven in that the analyst picks the bandwidth that they think should yield unbiased estimates.

The `rdd` package implements the Imbens-Kalyanaraman method to approach this problem. Imbens and Kalyanaraman advocate for optimizing the mean squared error using an algorithm that basically:

- Chooses an initial bandwidth and calculates the conditional expectation function and variance of y at the cutpoint
- Chooses a second initial bandwidth and do the same thing but calculate a second derivative of the CEF
- Add a regularization penalty

By iterating on these steps, we can eventually find the optimal bandwidth. Luckily, this is also implemented for us and we can just leave the `bw` argument blank by default to do this calculation:


```
rdd_model <- REstimate(# formula. outcome regressed on running variable x .
  # add a | c1 + c2... to add additional covariates
  myoutcomenext ~ difshare,
  # data
  data = elections,
  # specify cutpoint
  cutpoint = 0)
```

How did the Imbens-Kalyanaraman bandwidth estimate compare to our choice of .25?

Challenge

Another option is to use cross-validation. The basic procedure here is:

- Choose several values of bandwidths to search through then for each bandwidth value:
 - Split the data into v-folds
 - Estimate a RDD model using that bandwidth and calculate the MSE in each fold
 - Average the MSE across folds
- Select the bandwidth with the lowest MSE

See if you can implement these steps for cross-validation on your own! In the solutions, we make use of a few of the more advanced/latest tools in R like `predict()`, `vfold` from `tidymodels`, and `purrr`, a functional programming library that is part of the tidyverse. You may use these tools, or whatever else you like to attempt this challenge! Find an optimal bandwidth using this procedure, and report your average treatment effect.

Also note that the `REstimate()` function returns an object with class “RD”. See if you can extract the model object from it for calculating the MSE.

Solution to CV Bandwidth

1. The first step here is to create a function that we will use to estimate a regression discontinuity and calculate the Mean Squared Error.

```
calculate_rdd_and_bandwidth <- function(df_split, bw) {
  rdd_model <- REstimate(# formula. outcome regressed on running variable x
    # add a | c1 + c2... to add additional covariates
    myoutcomenext ~ difshare,
    # data
    data = df_split,
    # specify cutpoint
    cutpoint = 0,
    # specify bandwidth from Lee paper
    bw = bw,
    # return a model object
    model = TRUE)

  # Use predict() to get predictions from the model, compare them to the actual values of myoutcomenext
  mse_data <- data.frame(pred = predict(rdd_model$model[[1]]),
    actual = (df_split %>% filter(difshare >= -bw & difshare <= bw))$difshare)
  return(mean((mse_data$actual - mse_data$pred)^2))
}
```

2. Now we need to split our data. The `vfold()` function from `tidymodels` is similar to `train_test_split()` or `Kfold()` in Python's `sklearn`
3. Then we use the `map()` function from `purrr` to grab the data associated with each split using the "assessment" feature. We can then use `compose()` to prepare our rdd's for each split. For now we'll set a constant `bw` to test our function.
4. Then we use `map2()` to map our tidy functions and our custom rdd function to each split. **Note:** This operation returns `v` tibbles instead of a numeric vector, so we bind rows and then grab the mean

```
# V-fold split
elections_vfold <- vfold_cv(elections, v = 10, repeats = 1)

elections_vfold <- elections_vfold %>%
  mutate(df_split = map(splits, assessment))

tidy_rdd_model <- purrr::compose( # compose multiple functions
  broom::tidy, # convert lm objects into tidy tibbles
  calculate_rdd_and_bandwidth
)

bw = .25

tidied_models <- elections_vfold %>%
  mutate(rdd = map2(df_split, bw, tidy_rdd_model))

mean(bind_rows(tidied_models$rdd)$x)
```

```
## [1] 0.2068613
```

5. Now that we know this works for one value of `bw`, let's loop through 10 values of `bw` by incrementing from .01 to 1 in a for loop and performing the same operations. We'll return our average mean squared errors to a list and see which one was the lowest.

```
bw_seq <- seq(.01, 1, .05)

mses <- c()
counter = 1

for (bw in bw_seq){
  elections_vfold <- vfold_cv(elections, v = 10, repeats = 1)

  elections_vfold <- elections_vfold %>%
    mutate(df_split = map(splits, assessment))

  tidy_rdd_model <- purrr::compose( # compose multiple functions
    broom::tidy, # convert lm objects into tidy tibbles
    calculate_rdd_and_bandwidth
  )

  tidied_models <- elections_vfold %>%
    mutate(rdd = map2(df_split, bw, tidy_rdd_model))
```

```

mses[counter] <- mean(bind_rows(tidied_models$RDD)$x)
counter = counter + 1
}

```

Looks like .1913 was our lowest in this case, corresponding to our 11th bw, .51. Let's check our ATT for this run:

```

RDestimate(# formula. outcome regressed on running variable x
           # add a | c1 + c2... to add additional covariates
           myoutcomenext ~ difshare,
           # data
           data = elections,
           # specify cutpoint
           cutpoint = 0,
           # specify bandwidth from Lee paper
           bw = .51,
           # return a model object
           model = TRUE)

```

```

##
## Call:
## RDestimate(formula = myoutcomenext ~ difshare, data = elections,
##           cutpoint = 0, bw = 0.51, model = TRUE)
##
## Coefficients:
##           LATE      Half-BW    Double-BW
##           0.5876      0.5177      0.6474

```